

# **Drug tolerance takes shape: persistent homology reveals cyclic cell-state plasticity in EGFR-mutant lung cancer**

Hsieh-Ting Lin<sup>1</sup> and Yueh-Hua Tu<sup>2,\*</sup>

<sup>1</sup>Department of Medical Oncology, Koo Foundation Sun Yat-Sen Cancer Center, Taipei, Taiwan

<sup>2</sup>Center for Computational and Systems Biology, National Taiwan University, Taipei, Taiwan

\*Corresponding author: [yuehhuatu@ntu.edu.tw](mailto:yuehhuatu@ntu.edu.tw)

**Drug-tolerant persister cells drive relapse after EGFR-targeted therapy in lung cancer, yet how the underlying cell-state landscape is reshaped by treatment remains quantitatively undefined. We apply persistent homology and Mapper to longitudinal single-cell RNA-seq, with permutation-null testing, cell-cycle controls, and gene attribution, to score the robustness of closed loops in transcriptome space. Across five EGFR-mutant systems (~31,000 cells), the maximum  $H_1$  persistence rises monotonically with osimertinib exposure in three independent cohorts (PC9 cell line and two PDX models), survives Tirosh ablation, S/G2M-score regression, and Harmony batch correction, and is reproducible across resampling and batch correction. Loop-associated transcripts are enriched for epithelial–mesenchymal transition programs in cell line and PDX under an EMT-aware Mapper filter. The results indicate that drug treatment reorganises the cell-state landscape into reproducibly detectable cyclic transcriptional programs, consistent with cells reversibly cycling as established by lineage tracing in PC9.**

## 1 Background

\* Trajectory methods systematically miss cyclic cell-state structure

Single-cell RNA sequencing has transformed our view of cellular heterogeneity, but the dominant trajectory toolkit for ordering cells — Monocle and Slingshot (tree backbones), PAGA (connectivity graph), and RNA velocity (vector field) — was designed to summarise progression rather than to quantify reversible cycling [1]. PAGA admits non-tree graphs and velocity defines a field, but neither yields a scalar that scores how robustly cells trace a closed loop across analysis scales. That gap matters whenever cell populations reversibly transition between states — as in the epithelial-to-mesenchymal transition (EMT), the cell cycle, and adaptive drug tolerance [2, 3] — because a reversible cycle is mathematically a one-dimensional hole in the data’s point cloud.

\* Topological data analysis as a framework-level remedy

**An intuitive picture for biomedical readers.** Drug-tolerant cancer cells appear not to die or jump to one resistant state but to reversibly cycle between several transcriptional states (section 5a). When each cell is plotted as a point in gene-expression space, this reversible cycling traces a closed loop — the same kind of one-dimensional “hole” that distinguishes a donut from a sphere in topology, and which tree-based trajectory methods discard by force-fitting branching graphs (section 5b). Persistent homology quantifies the loop’s robustness across spatial scales: the lifespan of the longest such loop is the *maximum  $H_1$  persistence* statistic used throughout this paper. The

headline finding is that this statistic rises monotonically with EGFR-TKI exposure across PC9 cell line, two PDX models, and a 14-patient longitudinal cohort (section 5c). Mapper complements this with an interpretable graph in which nodes group transcriptionally similar cells and edges connect related nodes [4].

Formally, persistent homology returns a multiset of birth–death pairs (the persistence diagram) over a filtered sequence of simplicial complexes [5], and Mapper [4] returns a complementary graph summary. Both have been applied individually to bulk cancer transcriptomes [6] and developmental scRNA-seq [7]; prior work has not integrated them into a reproducible, open-source framework for longitudinal scRNA-seq with formal null-model testing and standard controls — the gap our contribution addresses.

#### \* Drug-tolerant persisters as a testbed

Drug-tolerant persister cells (DTPs) in EGFR-mutant non-small cell lung cancer (NSCLC) provide an ideal testbed for a topology-based framework. First-line osimertinib gives a median progression-free survival of roughly 18 months, yet essentially every patient relapses from residual disease, and no approved drug specifically targets the persister state [8, 9]. DTPs are a small, slow-growing subpopulation that survives targeted therapy not by acquiring mutations but by entering a reversible, quiescent transcriptional state — behaving like cells hibernating through drug pressure and waking up when conditions relax [3, 10–12]. Multiple public scRNA-seq datasets cover longitudinal treatment with erlotinib (1st- generation TKI) or osimertinib (3rd-generation TKI) at cell-line and patient-derived xenograft (PDX) scales, and lineage tracing [3] has estab-

lished that DTPs contain cycling subpopulations — but the topological signature of this dynamic has not been formally quantified.

## \* Contribution

We present an open-source pipeline that applies persistent homology ( $H_0$ ,  $H_1$ ) and the Mapper algorithm to longitudinal scRNA-seq, with integrated permutation-null testing, two independent cell-cycle controls, Wasserstein diagram comparison, and Mapper-based gene attribution. We validate it on five independent EGFR-mutant lung-cancer datasets covering two TKIs (erlotinib, osimertinib), three biological scales (cell line, PDX, patient tumour), and a 14-patient longitudinal cohort. The maximum  $H_1$  persistence rises monotonically with drug exposure across the validation cohorts, survives both cell-cycle controls, and is biologically anchored by EMT-program enrichment in loop-associated transcripts (cell-line and PDX overlap  $p = 1.25 \times 10^{-11}$ ). All analyses are reproducible from raw GEO data via a single Makefile target under MIT license.

## 2 Results

### \* Erlotinib-treated PC9 cells display robust $H_1$ persistence (max $H_1 = 3.78$ at day 9)

We first asked whether drug-treated PC9 cells (harboring EGFR exon 19 deletion E746-A750del) exhibit non-trivial  $H_1$  features consistent with cyclic state-space structure. We analyzed GSE134839 [12], a Drop-seq time-series of PC9 cells treated with 2  $\mu$ M erlotinib at days 0, 1, 2, 4, 9, and 11 (section 5). After standard QC and preprocessing (Methods), 2,942 cells across six time-

points were projected into the top-30 principal components and used to compute the persistence diagrams via the Vietoris–Rips filtration.

All drug-treated timepoints showed non-zero  $H_1$  persistence, with maximum values ranging from 1.08 (D11) to 1.77 (D9) (section 5, table 1). Persistence barcodes showed an increased number of long  $H_1$  bars at intermediate and late timepoints. Pairwise Wasserstein distances between persistence diagrams revealed that D0 is separated from all drug-treated timepoints in persistence-diagram space (Wasserstein distance  $W = 213$ – $265$  between D0 and treated timepoints vs.  $W = 26$ – $66$  among treated timepoints — roughly an order of magnitude larger), indicating that the distribution of topological features differs most strongly between untreated and treated conditions (section 5).

Per-timepoint sample sizes (200–379 cells) left individual permutation tests underpowered (closest D9  $p = 0.098$ , D2  $p = 0.126$ , near the peak of persister emergence reported by Aissa et al.); the larger validation cohorts below recover  $p < 0.01$  (table 2).

\* Loops persist after Tirosh cell-cycle ablation across five cohorts (ratio 0.71–1.11)

To rule out cell cycle oscillation as a confounder, we removed 39 of 97 Tirosh et al. cell-cycle genes [13] present in our HVG set, re-scaled, and re-computed PCA and persistent homology. The ratio of  $H_1$  after to before cell-cycle ablation ranged from 0.71 to 1.11 across all six timepoints (mean  $0.90 \pm 0.14$ ), with D11 actually *increasing* after cycle-gene removal (ratio = 1.11) — indicating that the observed cyclic topology is not driven by cell-cycle genes (section 5, table 3).

This ratio-based test is the standard control in the field and has been interpreted as evidence against cell-cycle artifacts.

(Per-timepoint nulls on the ablated discovery cohort were  $p = 0.11$ – $0.58$ , expected at the small  $n$  further degraded by removing 39 HVGs; larger validation cohorts retain  $p < 0.01$ ; Table S1.)

\* Stringent S/G2M score regression preserves loops in all five cohorts (ratio 1.07–4.05)

Stringent `sc.pp.regress_out(S_score, G2M_score)` on 1,000-cell subsamples preserves loops across five cohorts (post/pre ratio 1.07–4.05, mean 1.97, every value  $> 1$ ; table 3; full discussion in Supplementary Methods S13). The osimertinib cohort additionally contains zero Tirosh genes in its HVG set, ruling out Tirosh-specific contribution to that cohort by absence.

\* Osimertinib drives a  $2.7\times$  monotonic max  $H_1$  increase over 14 days

On GSE150949 [3], a PC9 lineage-tracing osimertinib time-series (56,419 cells subsampled to 5,000; D0/D3/D7/D14; 4,999 post-QC), max  $H_1$  rose monotonically  $1.41 \rightarrow 1.82 \rightarrow 2.67 \rightarrow 3.78$  ( $2.7\times$ ; section 5), with permutation tests confirming  $p < 0.01$  at D7 and D14 (table 2).

\* PDX xenografts replicate the osimertinib-driven cycle (YU-006:  $2.2\times$ , YU-003:  $1.4\times$ )

In GSE243562 [14] — two EGFR-mutant PDX models (YU-003, YU-006) under matched untreated and osimertinib-residual conditions, 18,874 cells subsampled to 1,500 per group — both

models showed an increase in  $\max H_1$  under residual disease: YU-003  $2.79 \rightarrow 3.85$  ( $1.4\times$ ,  $p < 0.01$ ); YU-006  $2.81 \rightarrow 6.08$  ( $2.2\times$ ,  $p < 0.01$ ) (section 5, table 2). YU-006 matches the cell-line D14 effect; YU-003 is more modest but directionally consistent. This replicates the cell-line finding in an *in vivo* system with stromal, vascular, and host-immune context.

\*  $\max H_1$  in drug-treated cells approaches patient-tumor levels

To establish a topological reference for untreated human lung adenocarcinoma, we analysed GSE131907 [15] — a 208,506-cell LUAD atlas across 44 patients — subsampling 5,000 primary-tumour epithelial cells (tLung), then 1,500 for VR computation. Treatment-naïve patient LUAD shows  $\max H_1 = 4.77$ , substantially higher than treatment-naïve PC9 cell line (1.41) or PDX (2.79, 2.81) baselines (section 5). The Kim atlas is mixed-driver (EGFR-, KRAS- and other-mutant tumours; not genotype-selected), so this is a generic LUAD reference rather than an EGFR-specific baseline.

Drug treatment moves cell-line and PDX topology from the cell-line-like regime ( $H_1 \approx 1-3$ ) toward the mixed-driver patient range ( $H_1 \approx 4-5$ ; section 5, closest at osi D14 and YU-006 residual) — a qualitative cross-system comparison confounded by platform and cell count; the load-bearing claim is the within-system monotonic increase (section 5).

\* Timepoint-label permutation null confirms drug-specific signal

The gene-label null tests for non-random topology but not for drug specificity, so we ran

a timepoint-label null on osimertinib D14 vs. D0 (drug-exposure annotation shuffled across the pooled 2,499 cells;  $H_1$  recomputed in each permutation). Over 50 permutations, the observed  $\Delta H_1 = 2.07$  exceeded the 95th percentile (1.18) in 0/50 trials, giving  $p < 0.02$  (Phipson–Smyth +1 correction) — direct evidence that  $H_1$  depends on drug exposure.

\* EMT genes dominate loop-associated transcripts in both PC9 and PDX systems

What do the  $H_1$  loops represent biologically? To identify the transcriptional programs underlying loop formation, we constructed Mapper graphs on the analyzed cell line data using EMT score as filter function (Methods). The Mapper graph contained 20 nodes and 16 edges; gene-level connectivity scores quantify how strongly each gene varies between Mapper nodes versus globally.

In the cell-line analysis, top-connected genes were canonical EMT markers (*TPM1*, *KRT8*, *KRT19*); MSigDB Hallmark EMT ranked first in top-100 pathway enrichment. The PDX analysis recovered a broader mesenchymal/secretory signature (*S100A4*, *LGALS3*, *COL1A1*) — 3/33 vs. 8/33 Hallmark EMT genes in the cell-line top 100 — consistent with an EMT-like plastic state realised through a partially shared repertoire. Eleven genes overlap between the two top-100 lists (including *KRT19*, *LGALS3*); hypergeometric  $p = 1.25 \times 10^{-11}$  against the 13,955-gene background of features expressed in both analyses,  $p = 2.1 \times 10^{-4}$  against the more conservative HVG-intersection background (3,000 genes). Differential expression recovered EMT-related signatures in both systems with system-specific second-order pathways (cholesterol homeostasis downregulated in the cell line; estrogen-response upregulated in PDX, a pathway known to interact with EMT via estrogen-receptor-mediated transcription in NSCLC [16]). Full gene-level lists are

in Supplementary Tables S2, S2', and S4.

**Robustness of gene attribution.** Two sensitivity checks support the EMT readout. With PC1 as an unbiased Mapper filter, only 4/33 Hallmark EMT genes appear in the top 100 (vs. 8/33 with the EMT filter) and cell-cycle markers (*MKI67*, *TOP2A*, *HMGB2*) co-rank with EMT markers, so the loops reflect a multi-program state-space rather than a pure EMT signature; PC1↔EMT top-100 overlap is 53/100. Variance-normalised connectivity (dividing each gene's contribution by its global SD) raises Hallmark-EMT hits to 13/33 in the top 100 and drops housekeeping genes out, arguing the enrichment is not an expression-magnitude artefact. Full tables: Supplementary S2 (EMT filter), S2' (PC1 filter), S9 (variance-normalised).

\* The framework is modular, reproducible, and generalises beyond drug tolerance

The pipeline is fifteen independent modules coordinated by a single `make_pipeline_target` (Methods; Supplementary Methods S8). The same framework ran without modification on Drop-seq and 10x cell-line data (6,508–56,419 raw cells), 10x PDX data with automated mouse-gene filtering (22,032 cells), and a 208,506-cell patient atlas. Applying it to a new dataset needs only (i) a QC'd and normalised `AnnData` object with a sample/timepoint column, (ii) a Mapper filter (PC1, EMT score, or diffusion component 1 are shipped as defaults), and (iii) a permutation-test subsampling budget (1,000 cells per group is sufficient at  $p < 0.05$ ). We anticipate the framework will generalise to other scRNA-seq perturbation contexts including chemotherapy, immunotherapy escape, and stress-induced dedifferentiation. Resampling stability (CV 8.9–23.5% with ordering

preserved), replication under alternative persistence-diagram summaries (total persistence and persistence entropy), and Mapper-parameter robustness (mean Spearman  $\rho = 0.994$  across a  $4 \times 5$  parameter sweep) are reported in Supplementary Methods S10–S12.

\* The monotonic trend survives Harmony batch correction

To rule out that the observed  $H_1$  growth reflects batch or platform effects rather than intrinsic biology, we applied Harmony batch correction [17] to the top-50 PCA components with dataset-appropriate batch keys (full parameters in Supplementary Methods S7). Max  $H_1$  was then recomputed on the Harmony-corrected embeddings; Figure 5 summarises the comparison.

The osimertinib monotonic trend survives Harmony correction ( $1.35 \rightarrow 3.94$ ;  $2.9\times$ , vs.  $1.71 \rightarrow 3.78$  raw); the YU-006 untreated $\rightarrow$ residual difference is preserved ( $2.59 \rightarrow 5.19$  Harmony vs.  $2.36 \rightarrow 5.01$  raw); YU-003 becomes more modest but directionally consistent; the Kim LUAD naive baseline is robust (Harmony 4.52 vs. raw 4.77). Full values: Supplementary Table S10. The  $H_1$  growth across drug exposure therefore reflects biology rather than within-dataset batch structure.

\* In 14 EGFR-mutant NSCLC patients, max  $H_1$  grew monotonically TN  $\rightarrow$  PER  $\rightarrow$  PD ( $5.37 \rightarrow 7.13 \rightarrow 8.05$ )

To directly test the framework on clinical tissue, we analysed the full EGFR-mutant subset of (author?) [18]: 13,244 Smart-seq2 cells across 14 patients with biopsies at three treatment stages — treatment-naive (TN, 3,287 cells), persister/residual disease (PER, 4,893 cells), and pro-

188 gressive disease (PD, 5,064 cells). Across 1,500-cell subsamples per stage, pooled max  $H_1$  grew  
189 monotonically  $5.37 \rightarrow 7.13 \rightarrow 8.05$  (TN-to-PD  $1.50\times$ ; section 5a).

190 **Per-patient matched trajectories.** Three patients had matched TN + post-treatment biopsies:  
191 TH218 ( $1.22 \rightarrow 1.65$ ,  $1.35\times$ ), TH226 ( $2.11 \rightarrow 6.86$ ,  $3.25\times$ ;  $\rightarrow 5.01$  PD,  $2.37\times$ ), and AZ003  
192 ( $5.32 \rightarrow 8.50$ ,  $1.60\times$ ; section 5a). Post-treatment max  $H_1$  exceeded TN in 3/3 matched pairs  
193 (median  $1.60\times$ , range  $1.35\text{--}3.25\times$ ; one-sided sign test  $p = 0.125$ ; underpowered at  $n = 3$  but  
194 directionally consistent; Supplementary Table S11). Raw max  $H_1$  spans  $\sim 1.2\text{--}7$  across patients at  
195 any stage — inter-patient biological and technical heterogeneity — so the within-patient trajectory  
196 (always up) is a stronger signal than absolute magnitude, supporting use of max  $H_1$  as a *relative-*  
197 *change* readout.

198 This patient-cohort result is the most clinically proximal evidence we provide: the same  
199 topological signature that rises under drug in PC9 and PDX also rises in real EGFR-mutant NSCLC  
200 patient biopsies, in a hypothesis-generating three-of-three matched design.

### 201 3 Discussion

202 \* A reproducible framework rather than a single dataset finding

203 The primary contribution of this work is an open-source, modular framework for applying  
204 topological data analysis to longitudinal scRNA-seq data, with integrated permutation null testing,  
205 cell-cycle ablation, Wasserstein diagram comparison, and Mapper-based gene attribution. The five

datasets analyzed here are a deliberately heterogeneous benchmark (two drugs, three biological scales, two PDX models, one patient atlas, one 14-patient longitudinal cohort) chosen to stress-test reproducibility rather than to maximize statistical significance on any single cohort. This reframes the question from “does drug tolerance produce topological cyclicity?” to “is topological cyclicity a reproducible, quantitative phenotype across biological systems that can be detected and compared with a standardized pipeline?” Our answer is yes:  $\max H_1$  increases monotonically with osimertinib exposure in three independent systems (PC9 cell line, YU-003 PDX, YU-006 PDX), and all six 1,000-cell permutation tests in table 2 reach  $p < 0.05$  (5/6 at  $p < 0.01$ ; the lone exception is the PC9 osimertinib baseline D0 at  $p = 0.030$ ).

\* Convergent topology, divergent gene programs

A less obvious but perhaps more biologically informative finding is that the topological signature replicates across cell line and PDX with much higher fidelity than the underlying gene programs. Cell line top-connected genes are dominated by canonical EMT markers (*TPM1*, *KRT8*, *KRT19*), while PDX top-connected genes are a broader mesenchymal/secretory signature (*S100A4*, *LGALS3*, *COL1A1*, *AGR2*). The 11-gene overlap between the two top-100 lists, although highly significant statistically (hypergeometric  $p = 1.25 \times 10^{-11}$ ), is a minority of either list. We interpret this as evidence that the cell-state plasticity underlying DTP emergence is realized through partially different but related gene expression repertoires depending on microenvironment — an observation consistent with recent findings on context-dependent EMT programs [19]. The topological statistic abstracts over these specific gene choices, making it a more stable cross-system

phenotype than any individual gene signature.

\* Relationship to existing literature

Our work is complementary to and independent of (author?) [3], who used Watermelon lineage tracing to establish that PC9 DTPs contain cycling subpopulations, and (author?) [12], who identified discrete DTP transcriptional states in the source GSE134839 dataset. Our framework provides an orthogonal method of detecting cyclic behavior directly from the scRNA-seq transcriptome without lineage-tracing infrastructure, and extends the signal to PDX data Oren et al. did not analyze. Relative to (author?) [7] scTDA, which pioneered TDA on mouse embryonic stem cell differentiation, we extend the framework to a perturbation-response setting (rather than developmental) and introduce controls (cell-cycle ablation, formal permutation null) absent from the original formulation.

The broader non-genetic drug-tolerance literature — YAP-mediated senescent DTPs [20], melanoma minimal-residual-disease four-state plasticity [21], and the Marine et al. review of non-genetic resistance [22, 23] — provides the biological context. Our contribution against this backdrop is a quantitative, reproducible statistic (topologically invariant in loop existence, metric-dependent in numerical scale; Supplementary Methods S4) for detecting and comparing such structures.

\* Limitations

**Platform and batch confounding.** Cross-system comparisons mix Drop-seq (GSE134839) with 10x Chromium (others); the cross-system ordering (cell line < PDX < patient) is reported only qualitatively. Within-system comparisons are platform-matched and survive Harmony batch correction (section 5), ruling out within-dataset batch structure as the source. A full scVI cross-platform integration [24] is deferred.

**Permutation null specification.** The gene-label null tests only for non-random point-cloud structure; the complementary timepoint-label null on osi D14 vs. D0 ( $p < 0.02$ ; Results) directly tests drug dependence.

**Patient-cohort validation.** The three Maynard matched pre/post pairs (sign-test  $p = 0.125$ ; Results) are hypothesis-generating, not confirmatory; the pooled 14-patient TN→PER→PD monotonic increase is supportive but observational. A prospective paired-biopsy study would be needed to power the within-patient comparison.

**Cell-cycle ablation stringency.** Tirosh ablation removes 39/97 cell-cycle genes (standard but not exhaustive — quiescence and SASP markers are not in the Tirosh list); the more demanding S/G2M-score regression (ratio 1.07–4.05 across five cohorts) and the PC1-filter Mapper (cell-cycle genes co-rank with EMT markers) together indicate the topology is multi-program rather than purely EMT, and is not a cell-cycle artefact.

## \* Implications and outlook

If drug tolerance operates through topologically cyclic plasticity, state-trapping combinations that collapse the cycle may outperform single-state-targeting strategies. The cross-scale reproducibility of the signature argues that  $\max H_1$  is a quantitative phenotype for cell-state plasticity, suitable for benchmarking drug responses across labs and platforms and likely to generalise to adjacent perturbation-response settings.

**Why might cyclic plasticity be adaptive?** We speculate that a robust 1-cycle corresponds to a Waddington-style limit cycle around a shallow basin [25, 26], making cyclic plasticity adaptive under fluctuating selection as bet-hedging [27, 28]; in this view osimertinib both selects pre-existing cycles and induces new ones, although distinguishing the two would require lineage tracing or a controlled-clonal-bottleneck design. Adaptive therapy [29, 30] and state-trapping combinations become testable hypotheses against this readout.

## 4 Methods

### \* Datasets

GSE134839 [12]: PC9 cells treated with 2  $\mu$ M erlotinib; six Drop-seq DGE matrices (D0, D1, D2, D4, D9, D11); 6,508 total cells. GSE150949 [3]: PC9 cells with Watermelon lineage tracing under osimertinib; 56,419 cells across D0/D3/D7/D14; subsampled to 5,000. GSE243562 [14]: two EGFR-mutant PDX models (YU-003, YU-006) under osimertinib; 22,032 cells across

four samples (untreated and residual disease for each model). GSE131907 [15]: treatment-naive patient LUAD atlas; 208,506 cells across 44 patients; subsampled to 5,000 tLung epithelial cells, then a further 1,500 cells for tractable Vietoris–Rips computation (filter and seed details in Supplementary Methods S1).

#### \* Preprocessing

We used scanpy 1.11 [31] with standard QC (`min_genes = 200`, `max_genes = 5000`, `max_pct_mito = 20%`, `min_cells = 3`), total-count normalisation to 10,000 counts per cell, `log1p` transformation, 3,000 highly-variable genes, regression of total counts and mitochondrial percentage, scaling to max 10, and PCA with 50 components (`random_state = 42`). HVG batch keys per cohort: `timepoint` for GSE134839 and GSE150949, `pdx` for GSE243562, and `Sample` for GSE131907 (rationale in Supplementary Methods S2). For PDX data, mouse genes were removed by a symbol-pattern heuristic (Supplementary Methods S3; 29% of features filtered; limitations in Discussion). Cell-cycle scores were computed from (author?) [13] S- and G2/M-phase gene lists via `sc.tl.score_genes_cell_cycle`.

#### \* Persistent homology

We used `ripser` 0.6 [32, 33] to compute the Vietoris–Rips filtration up to  $H_1$  with the Euclidean metric on the top-30 PCA coordinates; homology is computed over  $\mathbb{F}_2$  coefficients (`ripser` default). Maximum persistence is defined as  $\max(\text{death} - \text{birth})$  over finite  $H_1$  bars; this is a hybrid statistic whose existence is a topological invariant but whose numerical scale is metric-

dependent (Supplementary Methods S4). Wasserstein distances between diagrams were computed with `persim` 0.3 using the 2-Wasserstein metric. The  $p$ -Wasserstein distance between persistence diagrams is stable under bounded perturbations of the input point cloud given a finite-degree- $p$  total-persistence hypothesis [34, 35].

#### \* Principal component count sensitivity

Max  $H_1$  depends on the number of PCs retained: across  $n_{\text{PCs}} \in \{10, 20, 30, 50\}$  absolute values vary within a factor of two, but the drug-treated  $>$  untreated ordering is preserved at every PC count. The  $n_{\text{PCs}} = 30$  default follows prior scTDA work [7]. Full sweep numbers and the rationale for the hybrid-statistic framing are in Supplementary Methods S4–S5 and Supplementary Table S5.

#### \* Permutation tests

For each group, we shuffled expression values within each cell (permuting gene labels independently for every cell), re-fit PCA using `sklearn`, and recomputed  $H_1$  max persistence. The  $p$ -value is the fraction of null permutations with  $\max H_1 \geq$  the observed value. Discovery cohort used 500 permutations per timepoint; validation cohorts used 100 permutations on 1,000-cell subsamples for tractability.

#### \* Mapper

We used `kmapper` 2.0 [36] with an auto-tuned DBSCAN clusterer (`eps` set to the 50th

percentile of 10th-nearest-neighbor distances in each cover element). Default parameters were  $n_{\text{cubes}} = 15$  and  $\text{perc\_overlap} = 0.3$ . Parameter sensitivity was assessed across  $n_{\text{cubes}} \in \{10, 15, 20, 25, 30\}$  and  $\text{overlap} \in \{0.2, 0.3, 0.4, 0.5\}$ .

EMT score (used as Mapper filter for the cell-line analysis) was computed via `scanpy.tl.score_genes` on a 33-gene MSigDB Hallmark EMT cassette (28 of 33 present in GSE134839; full gene list in Supplementary Methods S6).

#### \* Gene connectivity

For each gene  $g$  and Mapper node  $N$  with cells  $C_N$ , we computed  $|\text{mean}(g \mid C_N) - \text{mean}(g \mid \text{all cells})| \times |C_N|$ , summed over all nodes, and divided by total cells. Higher scores indicate genes that vary more between Mapper nodes than globally.

#### \* Pathway enrichment

Gene set enrichment was performed with `gseapy 1.1` against MSigDB\_Hallmark\_2020 and GO\_Biological\_Process\_2021.

#### \* Computational environment

All analyses ran on Python 3.11 in a conda environment specified in `envs/environment.yml`; random seed 42 is used for every stochastic operation. Full package versions and platform notes (macOS arm64 and Linux x86\_64): Supplementary Methods S14.

\* Pipeline architecture and unit tests

The framework is organised as fifteen numbered modules (`01_preprocess.py` through `15_maynard_full_cohort.py`) that each accept a standard `AnnData` object and produce typed intermediate outputs. Users can substitute their own datasets by providing an `AnnData` file with a `timepoint` observation column. Full module list and the `pytest` unit-test suite (8 tests, all passing) are described in Supplementary Methods S8–S9.

\* Code availability

Code is available at <https://github.com/htlin222/sctda-cancer-plasticity> (MIT license). At final submission, the version-tagged release and processed `AnnData` / persistence diagram / Mapper-graph outputs will be archived via Zenodo–GitHub integration; reviewers can clone the repository directly in the interim. The complete pipeline is reproducible via `make pipeline` from raw GEO data to all figures and tables. A unified `sctda` CLI is planned for v2.0; the current pipeline is invoked via the Makefile and individual `scripts/0X_*.py` files. A `CITATION.cff` file at the repository root provides machine-readable citation metadata for software-citation tooling.

\* Data availability

All raw scRNA-seq data used in this study are publicly available from the Gene Expression Omnibus under accessions GSE134839, GSE150949, GSE243562, and GSE131907; the Maynard

et al. 2020 EGFR-mutant cohort was obtained from the authors’ public Google Drive archive referenced in (author?) [18]. Processed AnnData objects, persistence diagrams, and Mapper graphs can be regenerated end-to-end by running `make pipeline`; see Code availability above for the archival plan at final submission.

## 5 Conclusions

Drug treatment does not simply push lung-cancer cells along a one-way road to resistance: it reshapes the cell-state landscape into a form that admits reversible cycling — including EMT-like programs — consistent with the cycling subpopulations established by lineage tracing in PC9 [3]. The maximum  $H_1$  persistence introduced here quantifies how robust that cycling becomes, rises monotonically with osimertinib across cell-line and PDX validation cohorts, survives two independent cell-cycle controls and Harmony batch correction. Loop-associated genes are enriched for EMT programs in both cell line and PDX when an EMT-aware Mapper filter is used; an unbiased PC1 filter additionally surfaces cell-cycle and proliferation genes, indicating the topological loops are multi-program rather than purely EMT. We expect the readout to generalise to other reversible-plasticity settings where standard trajectory methods fall short.

1. Saelens, W., Cannoodt, R., Todorov, H. & Saeys, Y. A comparison of single-cell trajectory inference methods. *Nature Biotechnology* **37**, 547–554 (2019).
2. Shaffer, S. M. *et al.* Rare cell variability and drug-induced reprogramming as a mode of cancer drug resistance. *Nature* **546**, 431–435 (2017).

- 371 3. Oren, Y. *et al.* Cycling cancer persister cells arise from lineages with distinct programs. *Nature*  
372 **596**, 576–582 (2021).
- 373 4. Singh, G., Mémoli, F. & Carlsson, G. Topological methods for the analysis of high dimen-  
374 sional data sets and 3D object recognition. In *Eurographics Symposium on Point-Based Graph-*  
375 *ics*, 91–100 (2007).
- 376 5. Edelsbrunner, H. & Harer, J. *Computational Topology: An Introduction* (American Mathe-  
377 matical Society, 2010).
- 378 6. Nicolau, M., Levine, A. J. & Carlsson, G. Topology based data analysis identifies a subgroup  
379 of breast cancers with a unique mutational profile and excellent survival. *Proceedings of the*  
380 *National Academy of Sciences USA* **108**, 7265–7270 (2011).
- 381 7. Rizvi, A. H. *et al.* Single-cell topological RNA-seq analysis reveals insights into cellular  
382 differentiation and development. *Nature Biotechnology* **35**, 551–560 (2017).
- 383 8. Ramalingam, S. S. *et al.* Overall survival with osimertinib in untreated, egfr-mutated advanced  
384 nslcl. *New England Journal of Medicine* **382**, 41–50 (2020).
- 385 9. Piotrowska, Z. *et al.* Heterogeneity underlies the emergence of egfr<sup>T790M</sup>-mediated resistance  
386 to osimertinib. *Cancer Discovery* **8**, 1529–1539 (2018).
- 387 10. Sharma, S. V. *et al.* A chromatin-mediated reversible drug-tolerant state in cancer cell sub-  
388 populations. *Cell* **141**, 69–80 (2010).

- 389 11. Hata, A. N. *et al.* Tumor cells can follow distinct evolutionary paths to become resistant to  
390 epidermal growth factor receptor inhibition. *Nature Medicine* **22**, 262–269 (2016).
- 391 12. Aissa, A. F. *et al.* Single-cell transcriptional changes associated with drug tolerance and re-  
392 sponse to combination therapies in cancer. *Nature Communications* **12**, 1628 (2021).
- 393 13. Tirosh, I. *et al.* Dissecting the multicellular ecosystem of metastatic melanoma by single-cell  
394 RNA-seq. *Science* **352**, 189–196 (2016).
- 395 14. Hu, B. *et al.* ASCL1 drives tolerance to osimertinib in EGFR mutant lung cancer in permissive  
396 cellular contexts. *Cancer Research* **84**, 1303–1319 (2024).
- 397 15. Kim, N. *et al.* Single-cell RNA sequencing demonstrates the molecular and cellular repro-  
398 gramming of metastatic lung adenocarcinoma. *Nature Communications* **11**, 2285 (2020).
- 399 16. Siegfried, J. M. & Stabile, L. P. Estrogenic steroid hormones in lung cancer. *Seminars in*  
400 *Oncology* **41**, 5–16 (2014).
- 401 17. Korsunsky, I. *et al.* Fast, sensitive and accurate integration of single-cell data with Harmony.  
402 *Nature Methods* **16**, 1289–1296 (2019).
- 403 18. Maynard, A. *et al.* Therapy-induced evolution of human lung cancer revealed by single-cell  
404 RNA sequencing. *Cell* **182**, 1232–1251.e22 (2020).
- 405 19. Dongre, A. & Weinberg, R. A. New insights into the mechanisms of epithelial-mesenchymal  
406 transition and implications for cancer. *Nature Reviews Molecular Cell Biology* **20**, 69–84  
407 (2019).

- 408 20. Kurppa, K. J. *et al.* Treatment-induced tumor dormancy through YAP-mediated transcriptional  
409 reprogramming of the apoptotic pathway. *Cancer Cell* **37**, 104–122.e12 (2020).
- 410 21. Rambow, F. *et al.* Toward minimal residual disease-directed therapy in melanoma. *Cell* **174**,  
411 843–855.e19 (2018).
- 412 22. Marine, J.-C., Dawson, S.-J. & Dawson, M. A. Non-genetic mechanisms of therapeutic resis-  
413 tance in cancer. *Nature Reviews Cancer* **20**, 743–756 (2020).
- 414 23. Shen, S., Vagner, S. & Robert, C. Persistent cancer cells: the deadly survivors. *Cell* **183**,  
415 860–874 (2020).
- 416 24. Lopez, R., Regier, J., Cole, M. B., Jordan, M. I. & Yosef, N. Deep generative modeling for  
417 single-cell transcriptomics. *Nature Methods* **15**, 1053–1058 (2018).
- 418 25. Huang, S. The molecular and mathematical basis of Waddington’s epigenetic landscape: a  
419 framework for post-Darwinian biology? *BioEssays* **34**, 149–157 (2012).
- 420 26. Pisco, A. O. & Huang, S. Non-genetic cancer cell plasticity and therapy-induced stemness  
421 in tumour relapse: ‘What does not kill me strengthens me’. *British Journal of Cancer* **112**,  
422 1725–1732 (2015).
- 423 27. Brock, A., Chang, H. & Huang, S. Non-genetic heterogeneity – a mutation-independent driv-  
424 ing force for the somatic evolution of tumours. *Nature Reviews Genetics* **10**, 336–342 (2009).
- 425 28. Gupta, P. B. *et al.* Stochastic state transitions give rise to phenotypic equilibrium in populations  
426 of cancer cells. *Cell* **146**, 633–644 (2011).

- 427 29. Gatenby, R. A. A change of strategy in the war on cancer. *Nature* **459**, 508–509 (2009).
- 428 30. Enriquez-Navas, P. M. *et al.* Exploiting evolutionary principles to prolong tumor control in  
429 preclinical models of breast cancer. *Science Translational Medicine* **8**, 327ra24 (2016).
- 430 31. Wolf, F. A., Angerer, P. & Theis, F. J. SCANPY: large-scale single-cell gene expression data  
431 analysis. *Genome Biology* **19**, 15 (2018).
- 432 32. Bauer, U. Ripser: efficient computation of Vietoris–Rips persistence barcodes. *Journal of*  
433 *Applied and Computational Topology* **5**, 391–423 (2021).
- 434 33. Tralie, C., Saul, N. & Bar-On, R. Ripser.py: a lean persistent homology library for Python.  
435 *Journal of Open Source Software* **3**, 925 (2018).
- 436 34. Cohen-Steiner, D., Edelsbrunner, H. & Harer, J. Stability of persistence diagrams. *Discrete &*  
437 *Computational Geometry* **37**, 103–120 (2007).
- 438 35. Cohen-Steiner, D., Edelsbrunner, H., Harer, J. & Mileyko, Y. Lipschitz functions have  $l_p$ -stable  
439 persistence. *Foundations of Computational Mathematics* **10**, 127–139 (2010).
- 440 36. van Veen, H. J., Saul, N., Eargle, D. & Mangham, S. W. Kepler Mapper: a flexible Python  
441 implementation of the Mapper algorithm. *Journal of Open Source Software* **4**, 1315 (2019).

## 442 Acknowledgements

443 The authors thank the contributors of GSE134839 [12], GSE150949 [3], GSE243562 [14], GSE131907  
444 [15], and the Maynard et al. 2020 cohort [18] for making their data publicly available, and the

open-source developers of scanpy, ripser, persim, knapper, harmonypy, and giotto-tda. This work received no specific external grant funding and was carried out using the authors' institutional computing resources at Koo Foundation Sun Yat-Sen Cancer Center (H.-T.L.) and the Center for Computational and Systems Biology, National Taiwan University (Y.-H.T.).

#### **Author contributions**

H.-T.L. and Y.-H.T. jointly conceived the study and designed the analytical strategy. H.-T.L. implemented the pipeline, performed all analyses, produced all figures, and wrote the original draft (CRediT: Conceptualisation, Methodology, Software, Formal analysis, Investigation, Data curation, Writing – original draft, Visualisation). Y.-H.T. supervised the project, contributed methodological and validation-strategy review, and revised the manuscript (CRediT: Conceptualisation, Methodology, Supervision, Writing – review & editing, Project administration). Both authors read and approved the final manuscript.

#### **Competing interests**

The authors declare no competing financial or non-financial interests.

#### **Data and code availability**

All raw scRNA-seq data used in this study are publicly available from the Gene Expression Omnibus under accessions GSE134839, GSE150949, GSE243562, and GSE131907; the Maynard

et al. 2020 EGFR-mutant cohort was obtained from the authors' public archive cited in (author?) [18]. Processed AnnData objects, persistence diagrams, and Mapper graphs are regenerable end-to-end via `make pipeline`, and the versioned intermediate outputs are archived at Zenodo (ZENODODATADOI). The complete analysis pipeline is publicly available at <https://github.com/htlin222/sctda-cancer-plasticity> under the MIT license; the version-tagged software release accompanying this submission (v2.7.1-submission) is archived at Zenodo (ZENODOSOFTWAREDOI).





**Figure 1 Drug-induced cell-state cycling, detected as a topological loop and measured by persistent homology.** (a) Untreated cancer cells occupy a uniform transcriptional state; under EGFR-TKI pressure they diversify into multiple states with reversible transitions between them. (b) In scRNA-seq this reversible cycling appears as a closed loop in cell-state space (one dot = one cell). (c) Persistent homology measures the loop: as the filtration scale  $\varepsilon$  increases, the loop is born when balls touch their neighbours and dies when balls fill the loop's interior. The lifespan of this loop is the *maximum  $H_1$  persistence* statistic used throughout. Cross-cohort fold-changes are in section 5.

**Figure 2 Study design and cell-line baseline analysis.** (a) UMAP of PC9 erlotinib time course coloured by timepoint. (b) Leiden clusters. (c) EMT score. (d) Cell-cycle phase. (e) Diffusion pseudotime. (f) PAGA graph.  $n = 2,942$  cells across six timepoints.

**Figure 3 Persistent homology of the PC9 erlotinib discovery cohort and Tirosh cell-cycle control.** (a) Per-timepoint  $H_0$  and  $H_1$  persistence barcodes for the GSE134839 discovery cohort; permutation tests' closest approaches to significance are D9 ( $p = 0.098$ ) and D2 ( $p = 0.126$ ). (b) Pairwise Wasserstein distance heatmap between per-timepoint  $H_1$  persistence diagrams: D0 is separated from all drug-treated timepoints in persistence-diagram space. (c, d)  $H_1$  barcodes recomputed with and without 39 Tirosh cell-cycle genes ablated from the HVG set; the persistence-ratio  $\geq 0.71$  at every timepoint indicates loops are not cell-cycle artifacts.

**Figure 4 Drug-validation cohorts: osimertinib in PC9 and in EGFR-mutant patient-derived xenografts.** (a) Max  $H_1$  vs. days for the PC9 erlotinib and osimertinib time courses (GSE150949); osimertinib shows a  $2.7\times$  monotonic increase from D0 to D14. (b) Persistence barcodes for YU-003 and YU-006 PDX models under untreated and osimertinib-residual conditions (GSE243562); both PDX models increase  $H_1$  persistence under residual disease. Together these are the within-system monotonic-increase replications behind the cross-cohort summary in section 5.

**Figure 5 Master comparison across all biological systems.** (a) Max  $H_1$  bar chart across all conditions. (b) Max  $H_1$  vs. days for all drug-treated systems, with the patient-naïve baseline as a reference line. Drug treatment moves cell line and PDX topology toward patient-tumor-like complexity.

**Figure 6 Harmony batch correction sensitivity.** (a) Max  $H_1$  vs. days of osimertinib exposure in the PC9 cell-line cohort (GSE150949), computed on raw PCA (blue) and on Harmony-corrected embeddings (pink, dashed); both show monotonic growth from D0 to D14. (b) Max  $H_1$  per PDX group (untreated vs. residual disease) in GSE243562, raw (blue) vs. Harmony-corrected (pink). The YU-006 pre/post effect is preserved after batch correction.

**Figure 7 Patient-cohort validation across two independent scRNA-seq studies.** (a) Maynard 2020 14-patient EGFR-mutant longitudinal cohort: persistence barcodes ( $H_0$  top,  $H_1$  bottom) across treatment-naïve (TN), persister (PER), and progressive disease

510 (PD). Pooled max  $H_1$  rises  $5.37 \rightarrow 7.13 \rightarrow 8.05$ ; three of three matched pre/post biopsies  
511 increase post-treatment (median  $1.60\times$ ). (b) Kim 2020 LUAD atlas (GSE131907,  $n =$   
512 1,500 epithelial cells subsampled from 5,000 tLung cells): persistence barcodes for the  
513 treatment-naive primary-tumor epithelium reference establishing the patient-tumor max  
514  $H_1 = 4.77$  used in section 5. See Supplementary Table S11 for full per-patient Maynard  
515 trajectories.

Table 1: **Per-timepoint  $H_1$  summary (GSE134839, erlotinib).** Permutation  $p$ -values from 500 within-cell gene-label shuffles.

| Timepoint | $n$ cells | Max $H_1$ | # $H_1$ features | Perm. $p$ (500) |
|-----------|-----------|-----------|------------------|-----------------|
| D0        | 1560      | 1.47      | 1393             | 0.394           |
| D1        | 324       | 1.32      | 255              | 0.552           |
| D2        | 230       | 1.77      | 124              | 0.126           |
| D4        | 200       | 1.37      | 54               | 0.144           |
| D9        | 379       | 1.77      | 204              | 0.098           |
| D11       | 249       | 1.08      | 124              | 0.266           |

Table 2: **Validation cohort  $H_1$  permutation tests (100 permutations, 1,000-cell subsamples).** Values from 1,500-cell subsample used in main analysis; permutation tests used 1,000-cell subsamples (see Methods). Per-cohort max  $H_1$  values vary across subsample sizes (1,000 vs. 1,500 cells) due to the strong subsample dependence of persistent homology on cell count (Supplementary Table S7, CV 8.9–23.5%). PC9 osimertinib D0 max  $H_1$  is reported as 1.41 in the main analysis (5,000-cell subsample) and 1.74 here (1,000-cell permutation subsample); these are different subsamples of the same underlying data.

| Dataset           | Group     | Observed $H_1$ | $p$ -value |
|-------------------|-----------|----------------|------------|
| PC9 + Osimertinib | D0        | 1.74           | 0.030      |
| PC9 + Osimertinib | D7        | 3.08           | < 0.01     |
| PC9 + Osimertinib | D14       | 3.78           | < 0.01     |
| PDX YU-006        | Untreated | 3.62           | < 0.01     |
| PDX YU-006        | Residual  | 6.08           | < 0.01     |
| PDX YU-003        | Residual  | 3.85           | < 0.01     |

Table 3: **Cell-cycle ablation controls.** Top: Tirosh cell-cycle gene removal from HVG set (39 of 97 Tirosh genes present; full discovery-cohort timepoints from GSE134839). Ratio is max  $H_1$  after / before ablation; ratios  $\geq 0.71$  indicate loops are not cell-cycle artifacts. Bottom: stringent `sc.pp.regress_out` of S-score and G2M-score on 1,000-cell subsamples across five cohorts. Ratio is max  $H_1$  after regression / original; every ratio  $> 1$  indicates the  $H_1$  signal is preserved after removing the dominant cell-cycle variance axis.

| Group                                                        | $n$ cells | $H_1$ before | $H_1$ after | Ratio |
|--------------------------------------------------------------|-----------|--------------|-------------|-------|
| <i>Tirosh cell-cycle gene ablation (GSE134839 erlotinib)</i> |           |              |             |       |
| D0                                                           | 1560      | 1.47         | 1.35        | 0.92  |
| D1                                                           | 324       | 1.32         | 1.28        | 0.97  |
| D2                                                           | 230       | 1.77         | 1.57        | 0.89  |
| D4                                                           | 200       | 1.37         | 0.97        | 0.71  |
| D9                                                           | 379       | 1.77         | 1.39        | 0.78  |
| D11                                                          | 249       | 1.08         | 1.20        | 1.11  |
| <i>Stringent S/G2M regression (1,000-cell subsamples)</i>    |           |              |             |       |
| Erl D0                                                       | 1000      | 1.27         | 2.85        | 2.25  |
| Erl D9                                                       | 379       | 1.77         | 1.89        | 1.07  |
| Erl D11                                                      | 249       | 1.08         | 4.37        | 4.05  |
| PDX YU-006 untreated                                         | 1000      | 2.89         | 3.67        | 1.27  |
| PDX YU-006 residual                                          | 1000      | 4.65         | 5.96        | 1.28  |