

Drug tolerance takes shape: persistent homology reveals cyclic cell-state plasticity in EGFR-mutant lung cancer

Hsieh-Ting Lin¹ (ORCID: [0009-0002-3974-4528](https://orcid.org/0009-0002-3974-4528)) and Yueh-Hua Tu^{2,*} (ORCID: [0000-0002-1554-1753](https://orcid.org/0000-0002-1554-1753))

¹Department of Medical Oncology, Koo Foundation Sun Yat-Sen Cancer Center, Taipei, Taiwan

²Center for Computational and Systems Biology, National Taiwan University, Taipei, Taiwan

*Corresponding author: yuehhuatu@ntu.edu.tw

Drug-tolerant persister cells drive relapse after EGFR-targeted therapy in lung cancer, yet how the underlying cell-state landscape is reshaped by treatment remains quantitatively undefined. We apply persistent homology and Mapper to longitudinal single-cell RNA-seq, with permutation-null testing, cell-cycle controls, and gene attribution, to score the robustness of closed loops in transcriptome space. Across five EGFR-mutant systems ($\sim 31,000$ cells), the maximum H_1 persistence rises monotonically with osimertinib exposure in three independent cohorts (PC9 cell line and two PDX models), survives Tirosh ablation, S/G2M-score regression, and Harmony batch correction, and is reproducible across resampling seeds. Loop-associated transcripts are enriched for epithelial–mesenchymal transition programs in cell line and PDX under an EMT-aware Mapper filter. The results indicate that drug treatment reorganises the cell-state landscape into reproducibly detectable inter-state cycling — distinct from the cell cycle and from earlier single-state quiescence accounts of drug tolerance — consistent with the cycling subpopulations established by lineage tracing in PC9.

1 Background

* Trajectory methods systematically miss cyclic cell-state structure

Single-cell RNA sequencing has transformed our view of cellular heterogeneity, but the dominant trajectory toolkit — Monocle and Slingshot (tree backbones), PAGA (connectivity graph), and RNA velocity (vector field) — was designed to summarise one-directional progression rather than to quantify reversible inter-state cycling, i.e. cells visiting and revisiting several transcriptional states [1]. PAGA admits non-tree graphs and velocity defines a field, but neither yields a scalar that scores how robustly cells trace a closed loop. It matters in EMT and adaptive drug tolerance [2, 3]. The canonical cell cycle is cyclic in only one direction per cell ($G1 \rightarrow S \rightarrow G2 \rightarrow M$) yet produces a snapshot H_1 signature indistinguishable from bidirectional inter-state cycling, so we explicitly control for it. Geometrically, inter-state cycling traces a closed loop — a one-dimensional hole that persistent homology detects; directionality is supplied separately by lineage tracing [3].

* Topological data analysis as a framework-level remedy

An intuitive picture for biomedical readers. Drug-tolerant cancer cells appear neither to die outright nor to jump into a single resistant state; instead, individual cells visit and revisit *several* transcriptional states over time (section 5a). When each cell is plotted as a point in gene-expression space, this inter-state cycling traces a closed loop — the same kind of one-dimensional “hole” that distinguishes a donut from a sphere in topology, and which tree-based trajectory methods discard

by force-fitting branching graphs (section 5b). Persistent homology quantifies the loop’s robustness across spatial scales: the lifespan of the longest such loop is the *maximum H_1 persistence* statistic used throughout this paper. The headline finding is that this statistic rises monotonically with EGFR-TKI exposure across PC9 cell line, two PDX models, and a 14-patient longitudinal cohort (section 5c). Mapper complements this with an interpretable graph in which nodes group transcriptionally similar cells and edges connect related nodes [4].

Formally, persistent homology returns a multiset of birth–death pairs (the persistence diagram) over a filtered sequence of simplicial complexes [5], and Mapper [4] returns a complementary graph summary. Both have been applied individually to bulk cancer transcriptomes [6] and developmental scRNA-seq [7]; prior work has not integrated them into a reproducible, open-source framework for longitudinal scRNA-seq with formal null-model testing and standard controls — the gap our contribution addresses.

* Drug-tolerant persisters as a testbed

Drug-tolerant persister cells (DTPs) in EGFR-mutant non-small cell lung cancer (NSCLC) provide an ideal testbed for a topology-based framework. First-line osimertinib gives a median progression-free survival of roughly 18 months, yet essentially every patient relapses from residual disease, and no approved drug specifically targets the persister state [8, 9]. DTPs are a small, slow-growing subpopulation that survives targeted therapy without acquiring mutations [10, 11]. Early work framed DTPs as a single quiescent state [10]; more recent lineage tracing in PC9 osimertinib refines this picture, showing DTPs span multiple transcriptional states and can revert to

drug-sensitive phenotypes [3, 12] — the inter-state cycling our framework quantifies, distinct from a single quiescence $\leftrightarrow$ activation switch. Public scRNA-seq datasets cover longitudinal erlotinib and osimertinib treatment at cell-line and PDX scales, but the topological signature has not been formally quantified.

* Contribution

We present an open-source pipeline that applies persistent homology (H_0 , H_1) and the Mapper algorithm to longitudinal scRNA-seq, with integrated permutation-null testing, two independent cell-cycle controls, Wasserstein diagram comparison, and Mapper-based gene attribution. We validate it on five independent EGFR-mutant lung-cancer datasets covering two TKIs (erlotinib, osimertinib), three biological scales (cell line, PDX, patient tumour), and a 14-patient longitudinal cohort. The maximum H_1 persistence rises monotonically with drug exposure across the validation cohorts, survives both cell-cycle controls, and is biologically anchored by EMT-program enrichment in loop-associated transcripts (cell-line and PDX overlap $p = 1.25 \times 10^{-11}$). All analyses are reproducible from raw GEO data via a single Makefile target under MIT license.

2 Results

* Erlotinib-treated PC9 cells display robust H_1 persistence (max $H_1 = 3.78$ at day 9)

We first asked whether drug-treated NSCLC PC9 cells exhibit non-trivial H_1 features consistent with cyclic state-space structure. We analysed an erlotinib time-series in PC9 (GSE134839

[12]; six timepoints D0–D11; section 5; Methods).

All drug-treated timepoints showed non-zero H_1 persistence, with maximum values ranging from 1.08 (D11) to 1.77 (D9) (section 5, table 1). Persistence barcodes showed an increased number of long H_1 bars at intermediate and late timepoints. Pairwise Wasserstein distances between persistence diagrams revealed that D0 is separated from all drug-treated timepoints in persistence-diagram space (Wasserstein distance $W = 213$ – 265 between D0 and treated timepoints vs. $W = 26$ – 66 among treated timepoints — roughly an order of magnitude larger), indicating that the distribution of topological features differs most strongly between untreated and treated conditions (section 5).

Per-timepoint sample sizes (200–379 cells) left individual permutation tests underpowered (closest D9 $p = 0.098$, D2 $p = 0.126$, near the peak of persister emergence reported by Aissa et al.); the larger validation cohorts below recover $p < 0.01$ (table 2).

* Loops persist after Tirosh cell-cycle ablation across five cohorts (ratio 0.71–1.11)

To rule out cell cycle oscillation as a confounder, we removed 39 of 97 Tirosh et al. cell-cycle genes [13] present in our HVG set, re-scaled, and re-computed PCA and persistent homology. The ratio of H_1 after to before cell-cycle ablation ranged from 0.71 to 1.11 across all six timepoints (mean 0.90 ± 0.14), with D11 actually *increasing* after cycle-gene removal (ratio = 1.11) — indicating that the observed cyclic topology is not driven by cell-cycle genes (section 5, table 3). This ratio-based test is the standard control in the field and has been interpreted as evidence against

97 cell-cycle artifacts.

98 (Per-timepoint nulls on the ablated discovery cohort were $p = 0.11$ – 0.58 , expected at the
99 small n further degraded by removing 39 HVGs; larger validation cohorts retain $p < 0.01$; Ta-
100 ble S1.)

101 * Stringent S/G2M score regression preserves loops in all five cohorts (ratio 1.07–4.05)

102 Stringent `sc.pp.regress_out(S_score, G2M_score)` on 1,000-cell subsamples
103 preserves loops across five cohorts (post/pre ratio 1.07–4.05, mean 1.97, every value > 1 ; table 3;
104 full discussion in Supplementary Methods S13). The osimertinib cohort additionally contains zero
105 Tirosh genes in its HVG set, ruling out Tirosh-specific contribution to that cohort by absence.

106 * Osimertinib drives a $2.7\times$ monotonic max H_1 increase over 14 days

107 In a PC9 osimertinib time-series with Watermelon lineage tracing (GSE150949 [3]; D0/D3/D7/D14;
108 Methods), max H_1 rose monotonically $1.41 \rightarrow 1.82 \rightarrow 2.67 \rightarrow 3.78$ ($2.7\times$; section 5), with per-
109 mutation tests confirming $p < 0.01$ at D7 and D14 (table 2).

110 * PDX xenografts replicate the osimertinib-driven cycle (YU-006: $2.2\times$, YU-003: $1.4\times$)

111 In two EGFR-mutant PDX models (YU-003, YU-006) under matched untreated and osimertinib-
112 residual conditions (GSE243562 [14]; Methods), both models showed an increase in max H_1 under
113 residual disease: YU-003 $2.79 \rightarrow 3.85$ ($1.4\times$, $p < 0.01$); YU-006 $2.81 \rightarrow 6.08$ ($2.2\times$, $p < 0.01$)

(section 5, table 2). YU-006 matches the cell-line D14 effect; YU-003 is more modest but directionally consistent. This replicates the cell-line finding in an *in vivo* system with stromal, vascular, and host-immune context.

* Max H_1 in drug-treated cells approaches patient-tumor levels

To establish a topological reference for untreated human lung adenocarcinoma, we analysed a treatment-naive patient LUAD atlas (GSE131907 [15]; Methods). Treatment-naive patient LUAD shows max $H_1 = 4.77$, substantially higher than treatment-naive PC9 cell line (1.41) or PDX (2.79, 2.81) baselines (section 5). The Kim atlas is mixed-driver (EGFR-, KRAS- and other-mutant tumours; not genotype-selected), so this is a generic LUAD reference rather than an EGFR-specific baseline.

Drug treatment moves cell-line and PDX topology from the cell-line-like regime ($H_1 \approx 1-3$) toward the mixed-driver patient range ($H_1 \approx 4-5$; section 5, closest at osi D14 and YU-006 residual) — a qualitative cross-system comparison confounded by platform and cell count; the load-bearing claim is the within-system monotonic increase (section 5).

* Timepoint-label permutation null confirms drug-specific signal

The gene-label null tests for non-random topology but not for drug specificity, so we ran a timepoint-label null on osimertinib D14 vs. D0 (drug-exposure annotation shuffled across the pooled 2,499 cells; H_1 recomputed in each permutation). Over 50 permutations, the observed

$\Delta H_1 = 2.07$ exceeded the 95th percentile (1.18) in 0/50 trials, giving $p < 0.02$ (Phipson–Smyth +1 correction) — direct evidence that H_1 depends on drug exposure.

* EMT genes dominate loop-associated transcripts in both PC9 and PDX systems

What do the H_1 loops represent biologically? To identify the transcriptional programs underlying loop formation, we constructed Mapper graphs on the analyzed cell line data using EMT score as filter function (Methods). The Mapper graph contained 20 nodes and 16 edges; gene-level connectivity scores quantify how strongly each gene varies between Mapper nodes versus globally.

In the cell-line analysis, top-connected genes were canonical EMT markers (*TPM1*, *KRT8*, *KRT19*); MSigDB Hallmark EMT ranked first in top-100 pathway enrichment. The PDX analysis recovered a broader mesenchymal/secretory signature (*S100A4*, *LGALS3*, *COL1A1*) — 3/33 vs. 8/33 Hallmark EMT genes in the cell-line top 100 — consistent with an EMT-like plastic state realised through a partially shared repertoire. Eleven genes overlap between the two top-100 lists (including *KRT19*, *LGALS3*); hypergeometric $p = 1.25 \times 10^{-11}$ against the 13,955-gene background of features expressed in both analyses, $p = 2.1 \times 10^{-4}$ against the more conservative HVG-intersection background (3,000 genes). Differential expression recovered EMT-related signatures in both systems with system-specific second-order pathways (cholesterol homeostasis downregulated in the cell line; estrogen-response upregulated in PDX, a pathway known to interact with EMT via estrogen-receptor-mediated transcription in NSCLC [16]). Full gene-level lists are in Supplementary Tables S2, S2', and S4.

Robustness of gene attribution. Two sensitivity checks support the EMT readout. With PC1 as an unbiased Mapper filter, only 4/33 Hallmark EMT genes appear in the top 100 (vs. 8/33 with the EMT filter) and cell-cycle markers (*MKI67*, *TOP2A*, *HMGB2*) co-rank with EMT markers, so the loops reflect a multi-program state-space rather than a pure EMT signature; PC1↔EMT top-100 overlap is 53/100. Variance-normalised connectivity (dividing each gene’s contribution by its global SD) raises Hallmark-EMT hits to 13/33 in the top 100 and drops housekeeping genes out, arguing the enrichment is not an expression-magnitude artefact. Full tables: Supplementary S2 (EMT filter), S2’ (PC1 filter), S9 (variance-normalised).

* The framework is modular, reproducible, and generalises beyond drug tolerance

The pipeline is fifteen independent modules coordinated by a single `make_pipeline_target` (Methods; Supplementary Methods S8). The same framework ran without modification on Drop-seq and 10x cell-line data (6,508–56,419 raw cells), 10x PDX data with automated mouse-gene filtering (22,032 cells), and a 208,506-cell patient atlas. Applying it to a new dataset needs only (i) a QC’d and normalised `AnnData` object with a sample/timepoint column, (ii) a Mapper filter (PC1, EMT score, or diffusion component 1 are shipped as defaults), and (iii) a permutation-test subsampling budget (1,000 cells per group is sufficient at $p < 0.05$). We anticipate the framework will generalise to other scRNA-seq perturbation contexts including chemotherapy, immunotherapy escape, and stress-induced dedifferentiation. Resampling stability (CV 8.9–23.5% with ordering preserved), replication under alternative persistence-diagram summaries (total persistence and persistence entropy), and Mapper-parameter robustness (mean Spearman $\rho = 0.994$ across a 4×5

parameter sweep) are reported in Supplementary Methods S10–S12.

* The monotonic trend survives Harmony batch correction

To rule out that the observed H_1 growth reflects batch or platform effects rather than intrinsic biology, we applied Harmony batch correction [17] to the top-50 PCA components with dataset-appropriate batch keys (full parameters in Supplementary Methods S7). Max H_1 was then recomputed on the Harmony-corrected embeddings; Figure 5 summarises the comparison.

The osimertinib monotonic trend survives Harmony correction ($1.35 \rightarrow 3.94$; $2.9\times$, vs. $1.71 \rightarrow 3.78$ raw); the YU-006 untreated→residual difference is preserved ($2.59 \rightarrow 5.19$ Harmony vs. $2.36 \rightarrow 5.01$ raw); YU-003 becomes more modest but directionally consistent; the Kim LUAD naive baseline is robust (Harmony 4.52 vs. raw 4.77). Full values: Supplementary Table S10. The H_1 growth across drug exposure therefore reflects biology rather than within-dataset batch structure.

* In 14 EGFR-mutant NSCLC patients, max H_1 grew monotonically TN → PER → PD ($5.37 \rightarrow 7.13 \rightarrow 8.05$)

To directly test the framework on clinical tissue, we analysed the EGFR-mutant subset of (author?) [18] (14 patients with biopsies at treatment-naive, persister/residual, and progressive disease stages; Methods). Pooled max H_1 grew monotonically across the three stages: TN $5.37 \rightarrow$ PER $7.13 \rightarrow$ PD 8.05 (TN-to-PD $1.50\times$; section 5a).

Per-patient matched trajectories. Three patients had matched TN + post-treatment biopsies: TH218 ($1.22 \rightarrow 1.65$, $1.35\times$), TH226 ($2.11 \rightarrow 6.86$, $3.25\times$; $\rightarrow 5.01$ PD, $2.37\times$), and AZ003 ($5.32 \rightarrow 8.50$, $1.60\times$; section 5a). Post-treatment max H_1 exceeded TN in 3/3 matched pairs (median $1.60\times$, range 1.35 – $3.25\times$; one-sided sign test $p = 0.125$; underpowered at $n = 3$ but directionally consistent; Supplementary Table S11). Raw max H_1 spans ~ 1.2 – 7 across patients at any stage — inter-patient biological and technical heterogeneity — so the within-patient trajectory (always up) is a stronger signal than absolute magnitude, supporting use of max H_1 as a *relative-change* readout.

This patient-cohort result is the most clinically proximal evidence we provide: the same topological signature that rises under drug in PC9 and PDX also rises in real EGFR-mutant NSCLC patient biopsies, in a hypothesis-generating three-of-three matched design.

3 Discussion

* A reproducible framework rather than a single dataset finding

The primary contribution of this work is an open-source, modular framework for applying topological data analysis to longitudinal scRNA-seq data, with integrated permutation null testing, cell-cycle ablation, Wasserstein diagram comparison, and Mapper-based gene attribution. The five datasets here span two drugs, three biological scales, and a 14-patient cohort, chosen to stress-test reproducibility rather than to maximise single-cohort significance. This reframes the question from “does drug tolerance produce topological cyclicity?” to “is topological cyclicity a reproducible

cross-system phenotype?” Our answer is yes: $\max H_1$ increases monotonically with osimertinib exposure in three independent systems (PC9 cell line, YU-003 PDX, YU-006 PDX), and all six 1,000-cell permutation tests in table 2 reach $p < 0.05$ (5/6 at $p < 0.01$; the lone exception is the PC9 osimertinib baseline D0 at $p = 0.030$).

* Convergent topology, divergent gene programs

A less obvious but biologically informative finding is that the topological signature replicates across cell line and PDX with higher fidelity than the underlying gene programs. Cell-line top-connected genes are dominated by canonical EMT markers (*TPM1*, *KRT8*, *KRT19*); PDX top genes form a broader mesenchymal/secretory signature (*S100A4*, *LGALS3*, *COL1A1*, *AGR2*). The 11-gene overlap between the two top-100 lists is highly significant (hypergeometric $p = 1.25 \times 10^{-11}$) but a minority of either list. We interpret this as evidence that plasticity is realised through partially different gene repertoires depending on microenvironment, consistent with context-dependent EMT programs [19]. The topological statistic abstracts over these specific gene choices, making it a more stable cross-system phenotype than any individual gene signature.

* Relationship to existing literature

Our framework complements (author?) [3], who used Watermelon lineage tracing to establish PC9 DTP cycling subpopulations, and (author?) [12], who identified discrete DTP states in GSE134839. It detects cyclic behaviour directly from scRNA-seq without lineage-tracing infrastructure, and extends the signal to PDX data. Relative to (author?) [7] scTDA on mouse ES differ-

entiation, we extend the framework to perturbation-response and add cell-cycle / permutation-null controls absent from the original.

The broader non-genetic drug-tolerance literature — YAP-mediated senescent DTPs [20], melanoma minimal-residual-disease four-state plasticity [21], and the Marine et al. review of non-genetic resistance [22, 23] — provides the biological context. Our contribution against this backdrop is a quantitative, reproducible statistic (topologically invariant in loop existence, metric-dependent in numerical scale; Supplementary Methods S4) for detecting and comparing such structures.

* Limitations

Platform and batch confounding. Cross-system comparisons mix Drop-seq (GSE134839) with 10x Chromium (others); the cross-system ordering (cell line < PDX < patient) is reported only qualitatively. Within-system comparisons are platform-matched and survive Harmony batch correction (section 5), ruling out within-dataset batch structure as the source. A full scVI cross-platform integration [24] is deferred.

Permutation null specification. The gene-label null tests only for non-random point-cloud structure; the complementary timepoint-label null on osi D14 vs. D0 ($p < 0.02$; Results) directly tests drug dependence.

Patient-cohort validation. The three Maynard matched pre/post pairs (sign-test $p = 0.125$; Results) are hypothesis-generating, not confirmatory; the pooled 14-patient TN→PER→PD monotonic increase is supportive but observational. A prospective paired-biopsy study would be needed to power the within-patient comparison.

Cell-cycle ablation stringency. Tirosh ablation removes 39/97 cell-cycle genes (standard but not exhaustive — quiescence and SASP markers are not in the Tirosh list); the more demanding S/G2M-score regression (ratio 1.07–4.05 across five cohorts) and the PC1-filter Mapper (cell-cycle genes co-rank with EMT markers) together indicate the topology is multi-program rather than purely EMT, and is not a cell-cycle artefact.

* Implications and outlook

If drug tolerance operates through topologically cyclic plasticity, state-trapping combinations that collapse the cycle may outperform single-state-targeting strategies. The bidirectional nature of inter-state cycling further provides a conceptual basis for the partial clinical efficacy of TKI rechallenge after a drug holiday in EGFR-mutant NSCLC: cells that survived prior treatment can revert to drug-sensitive states. The cross-scale reproducibility of the signature argues that H_1 is a quantitative phenotype for cell-state plasticity, suitable for benchmarking drug responses across labs and likely to generalise to adjacent perturbation-response settings.

Why might cyclic plasticity be adaptive? We speculate that a robust 1-cycle corresponds to a Waddington-style limit cycle around a shallow basin [25, 26], making cyclic plasticity adaptive under fluctuating selection as bet-hedging [27, 28]; in this view osimertinib both selects pre-existing cycles and induces new ones, although distinguishing the two would require lineage tracing or a controlled-clonal-bottleneck design. Adaptive therapy [29, 30] and state-trapping combinations become testable hypotheses against this readout.

4 Methods

* Datasets

All cell models harbour activating *EGFR* mutations sensitising them to EGFR-TKIs; PC9 carries the canonical exon 19 deletion E746-A750del. Per-cohort details: GSE134839 [12]: PC9 erlotinib time-series; six Drop-seq DGE matrices at 2 μ M (D0, D1, D2, D4, D9, D11); 6,508 total cells, 2,942 post-QC. GSE150949 [3]: PC9 with Watermelon lineage tracing under osimertinib; 10x Chromium; 56,419 cells across D0/D3/D7/D14; subsampled to 5,000 (4,999 post-QC). GSE243562 [14]: two EGFR-mutant PDX models (YU-003, YU-006) under matched untreated and osimertinib-residual conditions; 10x Chromium; 22,032 cells (18,874 post-QC) across four samples, subsampled to 1,500 per group for VR computation. GSE131907 [15]: treatment-naive patient LUAD atlas (mixed driver genotypes); 10x Chromium; 208,506 cells across 44 patients; subsampled to 5,000 tLung epithelial cells, then 1,500 for tractable Vietoris–Rips computation (filter and seed details in Supplementary Methods S1). Maynard 2020 [18]: EGFR-mutant NSCLC

longitudinal cohort; Smart-seq2; 13,244 epithelial cells across 14 patients with biopsies at three treatment stages (treatment-naive TN, 3,287 cells; persister/residual disease PER, 4,893 cells; progressive disease PD, 5,064 cells); 1,500-cell subsamples per stage.

* Preprocessing

We used scanpy 1.11 [31] with standard QC (`min_genes = 200`, `max_genes = 5000`, `max_pct_mito = 20%`, `min_cells = 3`), total-count normalisation to 10,000 counts per cell, `log1p` transformation, 3,000 highly-variable genes, regression of total counts and mitochondrial percentage, scaling to max 10, and PCA with 50 components (`random_state = 42`). HVG batch keys per cohort: `timepoint` for GSE134839 and GSE150949, `pdx` for GSE243562, and `Sample` for GSE131907 (rationale in Supplementary Methods S2). For PDX data, mouse genes were removed by a symbol-pattern heuristic (Supplementary Methods S3; 29% of features filtered; limitations in Discussion). Cell-cycle scores were computed from (author?) [13] S- and G2/M-phase gene lists via `sc.tl.score_genes_cell_cycle`.

* Persistent homology

We used ripser 0.6 [32, 33] to compute the Vietoris–Rips filtration up to H_1 with the Euclidean metric on the top-30 PCA coordinates; homology is computed over $\mathbb{F}_2$ coefficients (riper default). Maximum persistence is defined as $\max(\text{death} - \text{birth})$ over finite H_1 bars; this is a hybrid statistic whose existence is a topological invariant but whose numerical scale is metric-dependent (Supplementary Methods S4). Wasserstein distances between diagrams were computed

with `persim` 0.3 using the 2-Wasserstein metric. The p -Wasserstein distance between persistence diagrams is stable under bounded perturbations of the input point cloud given a finite-degree- p total-persistence hypothesis [34, 35].

* Principal component count sensitivity

Max H_1 depends on the number of PCs retained: across $n_{\text{PCs}} \in \{10, 20, 30, 50\}$ absolute values vary within a factor of two, but the drug-treated $>$ untreated ordering is preserved at every PC count. The $n_{\text{PCs}} = 30$ default follows prior scTDA work [7]. Full sweep numbers and the rationale for the hybrid-statistic framing are in Supplementary Methods S4–S5 and Supplementary Table S5.

* Permutation tests

For each group, we shuffled expression values within each cell (permuting gene labels independently for every cell), re-fit PCA using `sklearn`, and recomputed H_1 max persistence. The p -value is the fraction of null permutations with $\max H_1 \geq$ the observed value. Discovery cohort used 500 permutations per timepoint; validation cohorts used 100 permutations on 1,000-cell subsamples for tractability.

* Mapper

We used `kmapper` 2.0 [36] with an auto-tuned DBSCAN clusterer (`eps` set to the 50th percentile of 10th-nearest-neighbor distances in each cover element). Default parameters were

316 $n_{\text{cubes}} = 15$ and $\text{perc_overlap} = 0.3$. Parameter sensitivity was assessed across $n_{\text{cubes}} \in \{10, 15, 20, 25, 30\}$
317 and $\text{overlap} \in \{0.2, 0.3, 0.4, 0.5\}$.

318 EMT score (used as Mapper filter for the cell-line analysis) was computed via `scanpy.tl.score_genes`
319 on a 33-gene MSigDB Hallmark EMT cassette (28 of 33 present in GSE134839; full gene list in
320 Supplementary Methods S6).

321 * Gene connectivity

322 For each gene g and Mapper node N with cells C_N , we computed $|\text{mean}(g \mid C_N) - \text{mean}(g \mid$
323 all cells)| $\times |C_N|$, summed over all nodes, and divided by total cells. Higher scores indicate genes
324 that vary more between Mapper nodes than globally.

325 * Pathway enrichment

326 Gene set enrichment was performed with `gseapy 1.1` against `MSigDB_Hallmark_2020`
327 and `GO_Biological_Process_2021`.

328 * Computational environment

329 All analyses ran on Python 3.11 in a conda environment specified in `envs/environment.yml`;
330 random seed 42 is used for every stochastic operation. Full package versions and platform notes
331 (macOS arm64 and Linux x86_64): Supplementary Methods S14.

332 * Pipeline architecture and unit tests

The framework is organised as fifteen numbered modules (`01_preprocess.py` through `15_maynard_full_cohort.py`) that each accept a standard `AnnData` object and produce typed intermediate outputs. Users can substitute their own datasets by providing an `AnnData` file with a `timepoint` observation column. Full module list and the `pytest` unit-test suite (8 tests, all passing) are described in Supplementary Methods S8–S9.

* Code availability

Code is available at <https://github.com/htlin222/sctda-cancer-plasticity> (MIT license). The version-tagged software release accompanying this submission is `v2.7.1-submission`, archived as a public GitHub release at <https://github.com/htlin222/sctda-cancer-plasticity/releases/tag/v2.7.1-submission>. The complete pipeline is reproducible via `make` pipeline from raw GEO data to all figures and tables. A unified `sctda` CLI is planned for `v2.0`; the current pipeline is invoked via the Makefile and individual `scripts/0X_*.py` files. A `CITATION.cff` file at the repository root provides machine-readable citation metadata for software-citation tooling.

* Data availability

All raw scRNA-seq data used in this study are publicly available from the Gene Expression Omnibus under accessions GSE134839, GSE150949, GSE243562, and GSE131907; the Maynard et al. 2020 EGFR-mutant cohort was obtained from the authors' public Google Drive archive referenced in (author?) [18]. Processed `AnnData` objects, persistence diagrams, and Mapper graphs

are regenerable end-to-end from the raw GEO data by running `make pipeline`; seeds and parameters are pinned in the repository so derived outputs are bit-stable across runs.

5 Conclusions

Drug treatment does not simply push lung-cancer cells along a one-way road to resistance: it reshapes the cell-state landscape into a form that admits inter-state cycling across EMT-like programs, consistent with the cycling subpopulations established by lineage tracing in PC9 [3]. The maximum H_1 persistence introduced here quantifies how robust that cycling becomes, rises monotonically with osimertinib across cell-line and PDX validation cohorts, survives two independent cell-cycle controls and Harmony batch correction. Loop-associated genes are enriched for EMT programs in both cell line and PDX when an EMT-aware Mapper filter is used; an unbiased PC1 filter additionally surfaces cell-cycle and proliferation genes, indicating the topological loops are multi-program rather than purely EMT. We expect the readout to generalise to other reversible-plasticity settings where standard trajectory methods fall short.

1. Saelens, W., Cannoodt, R., Todorov, H. & Saeys, Y. A comparison of single-cell trajectory inference methods. *Nature Biotechnology* **37**, 547–554 (2019).
2. Shaffer, S. M. *et al.* Rare cell variability and drug-induced reprogramming as a mode of cancer drug resistance. *Nature* **546**, 431–435 (2017).
3. Oren, Y. *et al.* Cycling cancer persister cells arise from lineages with distinct programs. *Nature* **596**, 576–582 (2021).

4. Singh, G., Mémoli, F. & Carlsson, G. Topological methods for the analysis of high dimensional data sets and 3D object recognition. In *Eurographics Symposium on Point-Based Graphics*, 91–100 (2007).
5. Edelsbrunner, H. & Harer, J. *Computational Topology: An Introduction* (American Mathematical Society, 2010).
6. Nicolau, M., Levine, A. J. & Carlsson, G. Topology based data analysis identifies a subgroup of breast cancers with a unique mutational profile and excellent survival. *Proceedings of the National Academy of Sciences USA* **108**, 7265–7270 (2011).
7. Rizvi, A. H. *et al.* Single-cell topological RNA-seq analysis reveals insights into cellular differentiation and development. *Nature Biotechnology* **35**, 551–560 (2017).
8. Ramalingam, S. S. *et al.* Overall survival with osimertinib in untreated, egfr-mutated advanced nslcl. *New England Journal of Medicine* **382**, 41–50 (2020).
9. Piotrowska, Z. *et al.* Heterogeneity underlies the emergence of egfr^{T790M}-mediated resistance to osimertinib. *Cancer Discovery* **8**, 1529–1539 (2018).
10. Sharma, S. V. *et al.* A chromatin-mediated reversible drug-tolerant state in cancer cell subpopulations. *Cell* **141**, 69–80 (2010).
11. Hata, A. N. *et al.* Tumor cells can follow distinct evolutionary paths to become resistant to epidermal growth factor receptor inhibition. *Nature Medicine* **22**, 262–269 (2016).

- 390 12. Aissa, A. F. *et al.* Single-cell transcriptional changes associated with drug tolerance and re-
391 sponse to combination therapies in cancer. *Nature Communications* **12**, 1628 (2021).
- 392 13. Tirosh, I. *et al.* Dissecting the multicellular ecosystem of metastatic melanoma by single-cell
393 RNA-seq. *Science* **352**, 189–196 (2016).
- 394 14. Hu, B. *et al.* ASCL1 drives tolerance to osimertinib in EGFR mutant lung cancer in permissive
395 cellular contexts. *Cancer Research* **84**, 1303–1319 (2024).
- 396 15. Kim, N. *et al.* Single-cell RNA sequencing demonstrates the molecular and cellular repro-
397 gramming of metastatic lung adenocarcinoma. *Nature Communications* **11**, 2285 (2020).
- 398 16. Siegfried, J. M. & Stabile, L. P. Estrogenic steroid hormones in lung cancer. *Seminars in*
399 *Oncology* **41**, 5–16 (2014).
- 400 17. Korsunsky, I. *et al.* Fast, sensitive and accurate integration of single-cell data with Harmony.
401 *Nature Methods* **16**, 1289–1296 (2019).
- 402 18. Maynard, A. *et al.* Therapy-induced evolution of human lung cancer revealed by single-cell
403 RNA sequencing. *Cell* **182**, 1232–1251.e22 (2020).
- 404 19. Dongre, A. & Weinberg, R. A. New insights into the mechanisms of epithelial-mesenchymal
405 transition and implications for cancer. *Nature Reviews Molecular Cell Biology* **20**, 69–84
406 (2019).
- 407 20. Kurppa, K. J. *et al.* Treatment-induced tumor dormancy through YAP-mediated transcriptional
408 reprogramming of the apoptotic pathway. *Cancer Cell* **37**, 104–122.e12 (2020).

- 409 21. Rambow, F. *et al.* Toward minimal residual disease-directed therapy in melanoma. *Cell* **174**,
410 843–855.e19 (2018).
- 411 22. Marine, J.-C., Dawson, S.-J. & Dawson, M. A. Non-genetic mechanisms of therapeutic resis-
412 tance in cancer. *Nature Reviews Cancer* **20**, 743–756 (2020).
- 413 23. Shen, S., Vagner, S. & Robert, C. Persistent cancer cells: the deadly survivors. *Cell* **183**,
414 860–874 (2020).
- 415 24. Lopez, R., Regier, J., Cole, M. B., Jordan, M. I. & Yosef, N. Deep generative modeling for
416 single-cell transcriptomics. *Nature Methods* **15**, 1053–1058 (2018).
- 417 25. Huang, S. The molecular and mathematical basis of Waddington’s epigenetic landscape: a
418 framework for post-Darwinian biology? *BioEssays* **34**, 149–157 (2012).
- 419 26. Pisco, A. O. & Huang, S. Non-genetic cancer cell plasticity and therapy-induced stemness
420 in tumour relapse: ‘What does not kill me strengthens me’. *British Journal of Cancer* **112**,
421 1725–1732 (2015).
- 422 27. Brock, A., Chang, H. & Huang, S. Non-genetic heterogeneity – a mutation-independent driv-
423 ing force for the somatic evolution of tumours. *Nature Reviews Genetics* **10**, 336–342 (2009).
- 424 28. Gupta, P. B. *et al.* Stochastic state transitions give rise to phenotypic equilibrium in populations
425 of cancer cells. *Cell* **146**, 633–644 (2011).
- 426 29. Gatenby, R. A. A change of strategy in the war on cancer. *Nature* **459**, 508–509 (2009).

- 427 30. Enriquez-Navas, P. M. *et al.* Exploiting evolutionary principles to prolong tumor control in
428 preclinical models of breast cancer. *Science Translational Medicine* **8**, 327ra24 (2016).
- 429 31. Wolf, F. A., Angerer, P. & Theis, F. J. SCANPY: large-scale single-cell gene expression data
430 analysis. *Genome Biology* **19**, 15 (2018).
- 431 32. Bauer, U. Ripser: efficient computation of Vietoris–Rips persistence barcodes. *Journal of*
432 *Applied and Computational Topology* **5**, 391–423 (2021).
- 433 33. Tralie, C., Saul, N. & Bar-On, R. Ripser.py: a lean persistent homology library for Python.
434 *Journal of Open Source Software* **3**, 925 (2018).
- 435 34. Cohen-Steiner, D., Edelsbrunner, H. & Harer, J. Stability of persistence diagrams. *Discrete &*
436 *Computational Geometry* **37**, 103–120 (2007).
- 437 35. Cohen-Steiner, D., Edelsbrunner, H., Harer, J. & Mileyko, Y. Lipschitz functions have l_p -stable
438 persistence. *Foundations of Computational Mathematics* **10**, 127–139 (2010).
- 439 36. van Veen, H. J., Saul, N., Eargle, D. & Mangham, S. W. Kepler Mapper: a flexible Python
440 implementation of the Mapper algorithm. *Journal of Open Source Software* **4**, 1315 (2019).

441 Acknowledgements

442 The authors thank the contributors of GSE134839 [12], GSE150949 [3], GSE243562 [14], GSE131907
443 [15], and the Maynard et al. 2020 cohort [18] for making their data publicly available, and the
444 open-source developers of scanpy, ripser, persim, knapper, harmonypy, and giotto-tda. This work

received no specific external grant funding and was carried out using the authors' institutional computing resources at Koo Foundation Sun Yat-Sen Cancer Center (H.-T.L.) and the Center for Computational and Systems Biology, National Taiwan University (Y.-H.T.).

Author contributions

H.-T.L. and Y.-H.T. jointly conceived the study and designed the analytical strategy. H.-T.L. implemented the pipeline, performed all analyses, produced all figures, and wrote the original draft (CRediT: Conceptualisation, Methodology, Software, Formal analysis, Investigation, Data curation, Writing – original draft, Visualisation). Y.-H.T. supervised the project, contributed methodological and validation-strategy review, and revised the manuscript (CRediT: Conceptualisation, Methodology, Supervision, Writing – review & editing, Project administration). Both authors read and approved the final manuscript.

Competing interests

The authors declare no competing financial or non-financial interests.

Data and code availability

All raw scRNA-seq data used in this study are publicly available from the Gene Expression Omnibus under accessions GSE134839, GSE150949, GSE243562, and GSE131907; the Maynard et al. 2020 EGFR-mutant cohort was obtained from the authors' public archive cited in

462 **(author?)** [18]. Processed AnnData objects, persistence diagrams, and Mapper graphs are re-
463 generable end-to-end from the raw GEO data via `make pipeline`; seeds and parameters are
464 pinned in the repository so outputs are bit-stable across runs. The complete analysis pipeline is
465 publicly available at <https://github.com/htlin222/sctda-cancer-plasticity>
466 under the MIT license. The version-tagged software release accompanying this submission is
467 `v2.7.1-submission` ([https://github.com/htlin222/sctda-cancer-plasticity/](https://github.com/htlin222/sctda-cancer-plasticity/releases/tag/v2.7.1-submission)
468 [releases/tag/v2.7.1-submission](https://github.com/htlin222/sctda-cancer-plasticity/releases/tag/v2.7.1-submission)), which is the authoritative software citation for the
469 present work.

Figure 1 Drug-induced cell-state cycling, detected as a topological loop and measured by persistent homology. (a) Untreated cancer cells occupy a uniform transcriptional state; under EGFR-TKI pressure they diversify into multiple states with reversible transitions between them. (b) In scRNA-seq this reversible cycling appears as a closed loop in cell-state space (one dot = one cell). (c) Persistent homology measures the loop: as the filtration scale ε increases, the loop is born when balls touch their neighbours and dies when balls fill the loop's interior. The lifespan of this loop is the *maximum H_1 persistence* statistic used throughout. Cross-cohort fold-changes are in section 5.

Figure 2 Study design and cell-line baseline analysis. (a) UMAP of PC9 erlotinib time course coloured by timepoint. (b) Leiden clusters. (c) EMT score. (d) Cell-cycle phase. (e) Diffusion pseudotime. (f) PAGA graph. $n = 2,942$ cells across six timepoints.

Figure 3 Persistent homology of the PC9 erlotinib discovery cohort and Tirosh cell-cycle control. (a) Per-timepoint H_0 and H_1 persistence barcodes for the GSE134839 discovery cohort; permutation tests' closest approaches to significance are D9 ($p = 0.098$) and D2 ($p = 0.126$). (b) Pairwise Wasserstein distance heatmap between per-timepoint H_1 persistence diagrams: D0 is separated from all drug-treated timepoints in persistence-diagram space. (c, d) H_1 barcodes recomputed with and without 39 Tirosh cell-cycle genes ablated from the HVG set; the persistence-ratio ≥ 0.71 at every timepoint indicates loops are not cell-cycle artifacts.

Figure 4 Drug-validation cohorts: osimertinib in PC9 and in EGFR-mutant patient-derived xenografts. (a) Max H_1 vs. days for the PC9 erlotinib and osimertinib time courses (GSE150949); osimertinib shows a $2.7\times$ monotonic increase from D0 to D14. (b) Persistence barcodes for YU-003 and YU-006 PDX models under untreated and osimertinib-residual conditions (GSE243562); both PDX models increase H_1 persistence under residual disease. Together these are the within-system monotonic-increase replications behind the cross-cohort summary in section 5.

Figure 5 Master comparison across all biological systems. (a) Max H_1 bar chart across all conditions. (b) Max H_1 vs. days for all drug-treated systems, with the patient-naive baseline as a reference line. Drug treatment moves cell line and PDX topology toward patient-tumor-like complexity.

Figure 6 Harmony batch correction sensitivity. (a) Max H_1 vs. days of osimertinib exposure in the PC9 cell-line cohort (GSE150949), computed on raw PCA (blue) and on Harmony-corrected embeddings (pink, dashed); both show monotonic growth from D0 to D14. (b) Max H_1 per PDX group (untreated vs. residual disease) in GSE243562, raw (blue) vs. Harmony-corrected (pink). The YU-006 pre/post effect is preserved after batch correction.

Figure 7 Patient-cohort validation across two independent scRNA-seq studies. (a) Maynard 2020 14-patient EGFR-mutant longitudinal cohort: persistence barcodes (H_0 top, H_1 bottom) across treatment-naive (TN), persister (PER), and progressive disease

511 (PD). Pooled max H_1 rises $5.37 \rightarrow 7.13 \rightarrow 8.05$; three of three matched pre/post biopsies
512 increase post-treatment (median $1.60\times$). (b) Kim 2020 LUAD atlas (GSE131907, $n =$
513 1,500 epithelial cells subsampled from 5,000 tLung cells): persistence barcodes for the
514 treatment-naive primary-tumor epithelium reference establishing the patient-tumor max
515 $H_1 = 4.77$ used in section 5. See Supplementary Table S11 for full per-patient Maynard
516 trajectories.

Table 1: **Per-timepoint H_1 summary (GSE134839, erlotinib).** Permutation p -values from 500 within-cell gene-label shuffles.

Timepoint	n cells	Max H_1	# H_1 features	Perm. p (500)
D0	1560	1.47	1393	0.394
D1	324	1.32	255	0.552
D2	230	1.77	124	0.126
D4	200	1.37	54	0.144
D9	379	1.77	204	0.098
D11	249	1.08	124	0.266

Table 2: **Validation cohort H_1 permutation tests (100 permutations, 1,000-cell subsamples).** Values from 1,500-cell subsample used in main analysis; permutation tests used 1,000-cell subsamples (see Methods). Per-cohort max H_1 values vary across subsample sizes (1,000 vs. 1,500 cells) due to the strong subsample dependence of persistent homology on cell count (Supplementary Table S7, CV 8.9–23.5%). PC9 osimertinib D0 max H_1 is reported as 1.41 in the main analysis (5,000-cell subsample) and 1.74 here (1,000-cell permutation subsample); these are different subsamples of the same underlying data.

Dataset	Group	Observed H_1	p -value
PC9 + Osimertinib	D0	1.74	0.030
PC9 + Osimertinib	D7	3.08	< 0.01
PC9 + Osimertinib	D14	3.78	< 0.01
PDX YU-006	Untreated	3.62	< 0.01
PDX YU-006	Residual	6.08	< 0.01
PDX YU-003	Residual	3.85	< 0.01

Table 3: **Cell-cycle ablation controls.** Top: Tirosh cell-cycle gene removal from HVG set (39 of 97 Tirosh genes present; full discovery-cohort timepoints from GSE134839). Ratio is max H_1 after / before ablation; ratios ≥ 0.71 indicate loops are not cell-cycle artifacts. Bottom: stringent `sc.pp.regress_out` of S-score and G2M-score on 1,000-cell subsamples across five cohorts. Ratio is max H_1 after regression / original; every ratio > 1 indicates the H_1 signal is preserved after removing the dominant cell-cycle variance axis.

Group	n cells	H_1 before	H_1 after	Ratio
<i>Tirosh cell-cycle gene ablation (GSE134839 erlotinib)</i>				
D0	1560	1.47	1.35	0.92
D1	324	1.32	1.28	0.97
D2	230	1.77	1.57	0.89
D4	200	1.37	0.97	0.71
D9	379	1.77	1.39	0.78
D11	249	1.08	1.20	1.11
<i>Stringent S/G2M regression (1,000-cell subsamples)</i>				
Erl D0	1000	1.27	2.85	2.25
Erl D9	379	1.77	1.89	1.07
Erl D11	249	1.08	4.37	4.05
PDX YU-006 untreated	1000	2.89	3.67	1.27
PDX YU-006 residual	1000	4.65	5.96	1.28